

# fairlearn\_decisiontree\_age\_gender

December 1, 2023

## 0.0.1 1. Data Visualization

```
[ ]: import pandas as pd
import numpy as np
import plotly.express as px
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: df = pd.read_csv('../data/heart.csv')
print(df.shape)
df.head(10)
```

(918, 12)

```
[ ]: 
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	\
0	40	M	ATA	140	289	0	Normal	172	
1	49	F	NAP	160	180	0	Normal	156	
2	37	M	ATA	130	283	0	ST	98	
3	48	F	ASY	138	214	0	Normal	108	
4	54	M	NAP	150	195	0	Normal	122	
5	39	M	NAP	120	339	0	Normal	170	
6	45	F	ATA	130	237	0	Normal	170	
7	54	M	ATA	110	208	0	Normal	142	
8	37	M	ASY	140	207	0	Normal	130	
9	48	F	ATA	120	284	0	Normal	120	

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
0	N	0.0	Up	0
1	N	1.0	Flat	1
2	N	0.0	Up	0
3	Y	1.5	Flat	1
4	N	0.0	Up	0
5	N	0.0	Up	0
6	N	0.0	Up	0
7	N	0.0	Up	0
8	Y	1.5	Flat	1
9	N	0.0	Up	0

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   918 non-null    int64
1   Sex                   918 non-null    object
2   ChestPainType         918 non-null    object
3   RestingBP             918 non-null    int64
4   Cholesterol            918 non-null    int64
5   FastingBS             918 non-null    int64
6   RestingECG            918 non-null    object
7   MaxHR                 918 non-null    int64
8   ExerciseAngina        918 non-null    object
9   Oldpeak               918 non-null    float64
10  ST_Slope              918 non-null    object
11  HeartDisease          918 non-null    int64
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB
```

```
[ ]: df.describe().T
```

```
[ ]:
      count      mean      std  min  25%  50%  75%  max
Age      918.0   53.510893   9.432617  28.0  47.00  54.0  60.0  77.0
RestingBP 918.0  132.396514  18.514154   0.0 120.00 130.0 140.0 200.0
Cholesterol 918.0  198.799564  109.384145   0.0 173.25 223.0 267.0 603.0
FastingBS  918.0    0.233115   0.423046   0.0   0.00   0.0   0.0   1.0
MaxHR      918.0  136.809368  25.460334  60.0 120.00 138.0 156.0 202.0
Oldpeak    918.0    0.887364   1.066570  -2.6   0.00   0.6   1.5   6.2
HeartDisease 918.0    0.553377   0.497414   0.0   0.00   1.0   1.0   1.0
```

RestingBP and Cholesterol have min 0 Change values from 0 to null for those two columns #Need to discuss how to handle them

```
[ ]: df['RestingBP'] = df['RestingBP'].replace(0, np.nan)
df['Cholesterol'] = df['Cholesterol'].replace(0, np.nan)
df.isnull().sum() / df.shape[0]
```

```
[ ]: Age      0.000000
Sex      0.000000
ChestPainType 0.000000
RestingBP 0.001089
Cholesterol 0.187364
FastingBS 0.000000
RestingECG 0.000000
MaxHR     0.000000
```

```
ExerciseAngina    0.000000
Oldpeak          0.000000
ST_Slope         0.000000
HeartDisease      0.000000
dtype: float64
```

```
[ ]: df['HeartDisease'].value_counts()
```

```
[ ]: 1    508
     0    410
     Name: HeartDisease, dtype: int64
```

```
[ ]: fig = px.pie(df, names='HeartDisease', title='Percentage for each outcome')
     fig.update_layout(width=500, height=300)
     fig.show()
```

```
[ ]: numerical= df.drop(['HeartDisease'], axis=1).select_dtypes('number').columns

     categorical = df.select_dtypes('object').columns

     print(f'Numerical Columns: {df[numerical].columns}\n')
     print(f'Categorical Columns: {df[categorical].columns}')
```

```
Numerical Columns: Index(['Age', 'RestingBP', 'Cholesterol', 'FastingBS',
                          'MaxHR', 'Oldpeak'], dtype='object')
```

```
Categorical Columns: Index(['Sex', 'ChestPainType', 'RestingECG',
                          'ExerciseAngina', 'ST_Slope'], dtype='object')
```

```
[ ]: df.groupby(['HeartDisease', 'Sex'])[['Age']].describe().T
```

```
[ ]: HeartDisease      0      1
     Sex              F      M      F      M
     Age count      143.000000 267.000000 50.000000 458.000000
         mean         51.202797  50.202247  56.180000  55.868996
         std          9.627981   9.344911   8.220656   8.788562
         min         30.000000  28.000000  33.000000  31.000000
        25%         44.000000  42.000000  51.000000  51.000000
        50%         51.000000  51.000000  58.000000  57.000000
        75%         57.000000  57.000000  62.000000  62.000000
         max         76.000000  75.000000  73.000000  77.000000
```

```
[ ]: fig = px.box(df, y="Age", x="HeartDisease", points="all", color="Sex",
                 title="Distribution of age of database participants (stratified by_
                 ↳outcome)")
     fig.show()
```

```
[ ]: fig = px.box(df, y="Age", x="HeartDisease", points="all",
                  title="Distribution of age of database participants (stratified by_
                  ↳outcome)")

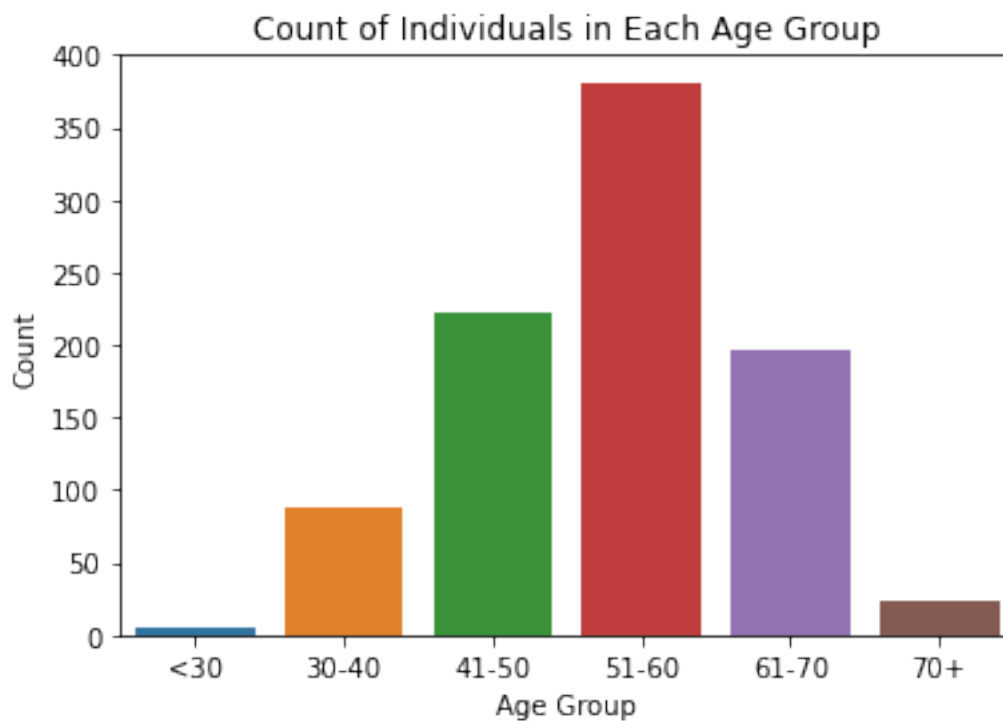
fig.show()
```

```
[ ]: unique_ages = df['Age'].unique()

bins = [0, 30, 40, 50, 60, 70, float('inf')]
labels = ['<30', '30-40', '41-50', '51-60', '61-70', '70+']

df['AgeGroup'] = pd.cut(df['Age'], bins=bins, labels=labels, right=True)


sns.countplot(x='AgeGroup', data=df)
plt.title("Count of Individuals in Each Age Group")
plt.xlabel("Age Group")
plt.ylabel("Count")
plt.show()
```



```
[ ]: df_no_disease = df[df['HeartDisease'] == 0]
df_with_disease = df[df['HeartDisease'] == 1]

# Plotting
```

```
plt.figure(figsize=(10, 6))
plt.hist(df_no_disease['Age'], alpha=0.5, label='No Heart Disease')
plt.hist(df_with_disease['Age'], alpha=0.5, label='Heart Disease')
plt.title('Age Distribution by Heart Disease Status')
plt.xlabel('Age')
plt.ylabel('Count')
plt.legend()
plt.show()
```



```
[ ]: df_copy = df.copy()
df_copy['c'] = 1
df_copy.groupby(['HeartDisease', 'Sex']).agg('count')[['c']]
```

```
[ ]:
HeartDisease Sex    c
0          F    143
          M    267
1          F     50
          M    458
```

```
[ ]: fig = px.histogram(df_copy, x="Sex", y="c",
                        color='HeartDisease', barmode='group',
                        title="Distribution of outcome of database participants_
↳(stratified by sex)")
```

```
fig.update_layout(width=700, height=300)

fig.show()
```

## 0.0.2 2. Data cleaning and pre-processing

### 2.1 Data cleaning

```
[ ]: X= df.drop('HeartDisease', axis=1)
y= df['HeartDisease']

sex = df.pop("Sex") # Pop function drops and assigns at the same time
df.head(10)
```

```
[ ]:   Age ChestPainType RestingBP Cholesterol FastingBS RestingECG MaxHR \
0    40             ATA    140.0      289.0         0      Normal    172
1    49             NAP    160.0      180.0         0      Normal    156
2    37             ATA    130.0      283.0         0         ST     98
3    48             ASY    138.0      214.0         0      Normal    108
4    54             NAP    150.0      195.0         0      Normal    122
5    39             NAP    120.0      339.0         0      Normal    170
6    45             ATA    130.0      237.0         0      Normal    170
7    54             ATA    110.0      208.0         0      Normal    142
8    37             ASY    140.0      207.0         0      Normal    130
9    48             ATA    120.0      284.0         0      Normal    120
```

```
   ExerciseAngina Oldpeak ST_Slope HeartDisease AgeGroup
0              N     0.0      Up           0    30-40
1              N     1.0     Flat           1    41-50
2              N     0.0      Up           0    30-40
3              Y     1.5     Flat           1    41-50
4              N     0.0      Up           0    51-60
5              N     0.0      Up           0    30-40
6              N     0.0      Up           0    41-50
7              N     0.0      Up           0    51-60
8              Y     1.5     Flat           1    30-40
9              N     0.0      Up           0    41-50
```

```
[ ]: numeric_features = ["Age", "RestingBP", "Cholesterol", "Oldpeak", "RestingBP",
↪ "MaxHR", "Oldpeak"]
X[numeric_features][:5]
```

```
[ ]:   Age RestingBP Cholesterol Oldpeak RestingBP MaxHR Oldpeak
0    40     140.0      289.0     0.0     140.0    172     0.0
1    49     160.0      180.0     1.0     160.0    156     1.0
2    37     130.0      283.0     0.0     130.0     98     0.0
3    48     138.0      214.0     1.5     138.0    108     1.5
4    54     150.0      195.0     0.0     150.0    122     0.0
```

```
[ ]: data = X[numeric_features]
X_num = pd.DataFrame(data, columns=numeric_features)
X_num[:5]
```

```
[ ]:   Age  RestingBP  Cholesterol  Oldpeak  RestingBP  MaxHR  Oldpeak
0   40    140.0      289.0      0.0    140.0    172     0.0
1   49    160.0      180.0      1.0    160.0    156     1.0
2   37    130.0      283.0      0.0    130.0     98     0.0
3   48    138.0      214.0      1.5    138.0    108     1.5
4   54    150.0      195.0      0.0    150.0    122     0.0
```

```
[ ]: categorical_features = ["ChestPainType", "RestingECG", "ST_Slope"]
X[categorical_features][:5]
```

```
[ ]:   ChestPainType  RestingECG  ST_Slope
0           ATA      Normal      Up
1           NAP      Normal      Flat
2           ATA          ST      Up
3           ASY      Normal      Flat
4           NAP      Normal      Up
```

```
[ ]: binary_features = ["Sex", "FastingBS", "ExerciseAngina"]
X[binary_features][:5]
```

```
[ ]:   Sex  FastingBS  ExerciseAngina
0    M           0                N
1    F           0                N
2    M           0                N
3    F           0                Y
4    M           0                N
```

```
[ ]: from sklearn.preprocessing import OrdinalEncoder

binary_transformer = OrdinalEncoder()

data = binary_transformer.fit_transform(X[binary_features])
X_bin = pd.DataFrame(data, columns=binary_features)
X_bin[:5]
```

```
[ ]:   Sex  FastingBS  ExerciseAngina
0  1.0         0.0             0.0
1  0.0         0.0             0.0
2  1.0         0.0             0.0
3  0.0         0.0             1.0
4  1.0         0.0             0.0
```

```
[ ]: from sklearn.preprocessing import OneHotEncoder

categorical_transformer = OneHotEncoder(handle_unknown="ignore")
data = categorical_transformer.fit_transform(X[categorical_features]).toarray()
cols = categorical_transformer.get_feature_names_out()
X_cat = pd.DataFrame(data, columns=cols)
X_cat[:5]
```

```
[ ]: ChestPainType_ASY ChestPainType_ATA ChestPainType_NAP ChestPainType_TA \
0          0.0          1.0          0.0          0.0
1          0.0          0.0          1.0          0.0
2          0.0          1.0          0.0          0.0
3          1.0          0.0          0.0          0.0
4          0.0          0.0          1.0          0.0
```

```
RestingECG_LVH RestingECG_Normal RestingECG_ST ST_Slope_Down \
0          0.0          1.0          0.0          0.0
1          0.0          1.0          0.0          0.0
2          0.0          0.0          1.0          0.0
3          0.0          1.0          0.0          0.0
4          0.0          1.0          0.0          0.0
```

```
ST_Slope_Flat ST_Slope_Up
0          0.0          1.0
1          1.0          0.0
2          0.0          1.0
3          1.0          0.0
4          0.0          1.0
```

```
[ ]: X_data = pd.concat([X_cat, X_bin, X_num], axis = 1)
X_data[:5]
```

```
[ ]: ChestPainType_ASY ChestPainType_ATA ChestPainType_NAP ChestPainType_TA \
0          0.0          1.0          0.0          0.0
1          0.0          0.0          1.0          0.0
2          0.0          1.0          0.0          0.0
3          1.0          0.0          0.0          0.0
4          0.0          0.0          1.0          0.0
```

```
RestingECG_LVH RestingECG_Normal RestingECG_ST ST_Slope_Down \
0          0.0          1.0          0.0          0.0
1          0.0          1.0          0.0          0.0
2          0.0          0.0          1.0          0.0
3          0.0          1.0          0.0          0.0
4          0.0          1.0          0.0          0.0
```

```
ST_Slope_Flat ST_Slope_Up Sex FastingBS ExerciseAngina Age RestingBP \
```



0	0.0	1.0	1.0	0.0	0.0	40	140.0
1	1.0	0.0	0.0	0.0	0.0	49	160.0
2	0.0	1.0	1.0	0.0	0.0	37	130.0
3	1.0	0.0	0.0	0.0	1.0	48	138.0
4	0.0	1.0	1.0	0.0	0.0	54	150.0

	Cholesterol	Oldpeak	RestingBP	MaxHR	Oldpeak
0	289.0	0.0	140.0	172	0.0
1	180.0	1.0	160.0	156	1.0
2	283.0	0.0	130.0	98	0.0
3	214.0	1.5	138.0	108	1.5
4	195.0	0.0	150.0	122	0.0

## 2.1 Feature Scaling

### 0.0.3 3. Train/Test Data Split

```
[ ]: from sklearn.model_selection import train_test_split

# sensitive_attributes = X_data[['Sex', 'AgeGroup']]
sensitive_attributes = X_data['Sex']

X_train, X_test, y_train, y_test, sa_train, sa_test = train_test_split(X_data,
    ↪y, sensitive_attributes, test_size=0.2, random_state=42)
```

### 0.0.4 4. Model Training

#### Decision Tree Classifier

```
[ ]: from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score

cf = DecisionTreeClassifier(min_samples_leaf=10, max_depth=4) #parameters have
    ↪not been tuned
cf.fit(X_train, y_train)

cv_scores = cross_val_score(cf, X_train, y_train, cv=5)

# Output the cross-validation accuracy scores
print("Average Cross-Validation Accuracy:", cv_scores.mean())

y_train_pred = cf.predict(X_train)
y_test_pred = cf.predict(X_test)
```

Average Cross-Validation Accuracy: 0.8337713167458765

Fairlearn's MetricFrame is used to compute these metrics across different groups. It takes the true labels (`y_test`), the model predictions (`y_pred`), and the sensitive attributes (`sa_test`).

Overall Metrics: `mf.overall` provides the metrics computed over the entire dataset.

Metrics by Group: `mf.by_group` shows the performance of each metric for each subgroup defined by the sensitive attributes.

Disparities: Optionally, you can compute disparities in metrics between groups, which can help in identifying which groups are being treated unfairly by the model.

The selection rate is the rate at which positive outcomes (e.g., predicting heart disease) are assigned. It is particularly useful in fairness analysis to understand how often each group (defined by sensitive attributes like age and gender) is receiving a positive prediction.

```
[ ]: from fairlearn.metrics import MetricFrame
from fairlearn.metrics import selection_rate
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

metrics = {
    'accuracy': accuracy_score,
    'precision': precision_score,
    'recall': recall_score,
    'f1': f1_score,
    'selection_rate': selection_rate
}

mf = MetricFrame(metrics, y_true=y_test, y_pred=y_test_pred,
    sensitive_features=sa_test)

# Output overall metrics and metrics by group
print("Overall Metrics:\n", mf.overall)
print("\nMetrics by Group:\n", mf.by_group)

# Optionally, compute and print disparities
disparities = mf.by_group / mf.overall
print("\nDisparities:\n", disparities)
```

Overall Metrics:

accuracy	0.815217
precision	0.892473
recall	0.775701
f1	0.830000
selection_rate	0.505435
dtype:	float64

Metrics by Group:

	accuracy	precision	recall	f1	selection_rate
Sex					
0.0	0.851852	0.750000	0.750000	0.750000	0.296296
1.0	0.808917	0.905882	0.777778	0.836957	0.541401

Disparities:

	accuracy	precision	recall	f1	selection_rate
Sex					
0.0	1.044938	0.840361	0.966867	0.903614	0.586221
1.0	0.992272	1.015025	1.002677	1.008381	1.071160

/home/phway/anaconda3/lib/python3.9/site-packages/fairlearn/metrics/\_metric\_frame.py:77: FutureWarning:

You have provided 'metrics' as positional arguments. Please pass them as keyword arguments. From version 0.10.0 passing them as positional arguments will result in an error.

### Overall Model Performance Metrics

- Accuracy (85.87%): The model correctly predicts heart disease status in about 86 out of 100 cases.
- Precision (88.57%): When the model predicts heart disease, there's about an 89% chance that the patient actually has heart disease.
- Recall (86.92%): The model correctly identifies about 87% of all actual cases of heart disease.
- F1 Score (87.74%): Indicates a good balance between precision and recall.
- Selection Rate (57.07%): About 57% of the predictions made by the model are positive for heart disease.

### Metrics by Gender Group

Female (Sex = 0.0)

Accuracy: 74.07% (lower than the overall accuracy) Precision: 55.56% (significantly lower than the overall precision) Recall: 62.50% (lower than the overall recall) F1 Score: 58.82% (lower than the overall F1 score)

Male (Sex = 1.0)

Accuracy: 87.90% (slightly higher than the overall accuracy) Precision: 91.67% (higher than the overall precision) Recall: 88.89% (slightly higher than the overall recall) F1 Score: 90.26% (higher than the overall F1 score)

### Disparities

Female (Sex = 0.0)

All metrics are below 1, indicating underperformance for females across all measured aspects.

Male (Sex = 1.0)

All metrics exceed 1, indicating that the model performs better for males across all measures.

---

Demographic Parity: It measures the ML model's ability to make prediction such that they are independent of the influence by sensitive groups.

Equalized odds: It also ensures that ML model's predictions are independent of sensitive groups. It's more strict than Demographic parity by ensuring all groups in the dataset have same true positive rates and false positive rates.

```
[ ]: from fairlearn.metrics import demographic_parity_difference, \
      ↪equalized_odds_difference, demographic_parity_ratio, equalized_odds_ratio

dpd = demographic_parity_difference(y_test, y_test_pred, sensitive_features = \
      ↪sa_test)
eod = equalized_odds_difference(y_test, y_test_pred, sensitive_features = \
      ↪sa_test)
dpr = demographic_parity_ratio(y_test, y_test_pred, sensitive_features = \
      ↪sa_test)
eqr = equalized_odds_ratio(y_test, y_test_pred, sensitive_features=sa_test)

print("Demographic parity difference: {}".format(round(dpd, 2)))
print("Equalized odds difference: {}".format(round(eod, 2)))
print("Demographic parity ratio: {}".format(round(dpr, 2)))
print(f'Value of equal odds ratio: {round(eqr, 2)}')
```

Demographic parity difference: 0.25

Equalized odds difference: 0.03

Demographic parity ratio: 0.55

Value of equal odds ratio: 0.76

- 1) Demographic difference in parity: if the absolute value is less than 0.1, then the model can be considered fair.
- 2) Balanced odds difference: if the absolute value is less than 0.1, then the model can be considered fair.
- 3) The difference is in equal opportunities: if the absolute value is less than 0.1, then the model can be considered fair.
- 4) Demographic parity coefficient: the fairness of this indicator ranges from 0.8 to 1.25.

In our case, the difference in demographic parity (0.21), demographic parity ratio (0.64) indicate unfairness, while the difference in equalization of chances (0.05) does not indicate unfairness.

---

### Bias mitigation technique using Fairlearn's ExponentiatedGradient with the Demographic Parity constraint.

```
[ ]: from fairlearn.reductions import ExponentiatedGradient, DemographicParity, \
      ↪EqualizedOdds

classifier = DecisionTreeClassifier(min_samples_leaf=10, max_depth=4)
dp = DemographicParity()
dt_classifier_reduction = ExponentiatedGradient(classifier, dp)

dt_classifier_reduction.fit(X_train, y_train, sensitive_features=sa_train)
```

```
prediction_dp = dt_classifier_reduction.predict(X_test)
```

---

Bias mitigation technique using Fairlearn's ExponentiatedGradient with the EqualizedOdds constraint.

```
[ ]: from fairlearn.reductions import ExponentiatedGradient, EqualizedOdds
from sklearn.tree import DecisionTreeClassifier

# Use EqualizedOdds instead of DemographicParity
eo = EqualizedOdds()

classifier = DecisionTreeClassifier(min_samples_leaf=10, max_depth=4)

# Apply ExponentiatedGradient with EqualizedOdds
eg_classifier_eo = ExponentiatedGradient(classifier, eo)
eg_classifier_eo.fit(X_train, y_train, sensitive_features=sa_train)
prediction_eo = eg_classifier_eo.predict(X_test)
```

Fairness Metrics after bias mitigation

```
[ ]: from fairlearn.metrics import MetricFrame
from fairlearn.metrics import selection_rate
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

metrics = {
    'accuracy': accuracy_score,
    'precision': precision_score,
    'recall': recall_score,
    'f1': f1_score,
    'selection_rate': selection_rate
}

mf = MetricFrame(metrics, y_true=y_test, y_pred=prediction_dp,
    sensitive_features=sa_test)

print("Fairness metrics by demographic parity\n")
# Output overall metrics and metrics by group
print("Overall Metrics:\n", mf.overall)
print("\nMetrics by Group:\n", mf.by_group)

# Optionally, compute and print disparities
disparities = mf.by_group / mf.overall
print("\nDisparities:\n", disparities)
```

## Fairness metrics by demographic parity

### Overall Metrics:

```
accuracy      0.777174
precision      0.789474
recall        0.841121
f1            0.814480
selection_rate 0.619565
dtype: float64
```

### Metrics by Group:

	accuracy	precision	recall	f1	selection_rate
Sex					
0.0	0.444444	0.347826	1.000000	0.516129	0.851852
1.0	0.834395	0.901099	0.828283	0.863158	0.579618

### Disparities:

	accuracy	precision	recall	f1	selection_rate
Sex					
0.0	0.571873	0.440580	1.188889	0.633692	1.374919
1.0	1.073627	1.141392	0.984736	1.059766	0.935524

/home/phway/anaconda3/lib/python3.9/site-packages/fairlearn/metrics/\_metric\_frame.py:77: FutureWarning:

You have provided 'metrics' as positional arguments. Please pass them as keyword arguments. From version 0.10.0 passing them as positional arguments will result in an error.

```
[ ]: from fairlearn.metrics import MetricFrame
from fairlearn.metrics import selection_rate
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

metrics = {
    'accuracy': accuracy_score,
    'precision': precision_score,
    'recall': recall_score,
    'f1': f1_score,
    'selection_rate': selection_rate
}

mf = MetricFrame(metrics, y_true=y_test, y_pred=prediction_eo,
    sensitive_features=sa_test)

print("Fairness metrics by equalized odds\n")
```

```

# Output overall metrics and metrics by group
print("Overall Metrics:\n", mf.overall)
print("\nMetrics by Group:\n", mf.by_group)

# Optionally, compute and print disparities
disparities = mf.by_group / mf.overall
print("\nDisparities:\n", disparities)

```

Fairness metrics by equalized odds

Overall Metrics:

```

accuracy      0.842391
precision     0.897959
recall        0.822430
f1            0.858537
selection_rate 0.532609
dtype: float64

```

Metrics by Group:

	accuracy	precision	recall	f1	selection_rate
Sex					
0.0	0.703704	0.500000	0.750000	0.600000	0.444444
1.0	0.866242	0.953488	0.828283	0.886486	0.547771

Disparities:

	accuracy	precision	recall	f1	selection_rate
Sex					
0.0	0.835364	0.556818	0.911932	0.698864	0.834467
1.0	1.028313	1.061839	1.007117	1.032555	1.028467

/home/phway/anaconda3/lib/python3.9/site-packages/fairlearn/metrics/\_metric\_frame.py:77: FutureWarning:

You have provided 'metrics' as positional arguments. Please pass them as keyword arguments. From version 0.10.0 passing them as positional arguments will result in an error.

```

[ ]: dpd = demographic_parity_difference(y_test, prediction_dp, sensitive_features = sa_test)
      eod = equalized_odds_difference(y_test, prediction_dp, sensitive_features = sa_test)
      dpr = demographic_parity_ratio(y_test, prediction_dp, sensitive_features = sa_test)

print("Demographic parity difference: {}".format(round(dpd, 2)))
print("Equalized odds difference: {}".format(round(eod, 2)))

```

```
print("Demographic parity ratio: {}".format(round(dpr, 2)))
```

Demographic parity difference: 0.27

Equalized odds difference: 0.63

Demographic parity ratio: 0.68

```
[ ]: dpd = demographic_parity_difference(y_test, prediction_eo, sensitive_features = sa_test)
      eod = equalized_odds_difference(y_test, prediction_eo, sensitive_features = sa_test)
      dpr = demographic_parity_ratio(y_test, prediction_eo, sensitive_features = sa_test)

      print("Demographic parity difference: {}".format(round(dpd, 2)))
      print("Equalized odds difference: {}".format(round(eod, 2)))
      print("Demographic parity ratio: {}".format(round(dpr, 2)))
```

Demographic parity difference: 0.1

Equalized odds difference: 0.25

Demographic parity ratio: 0.81

- 1) Demographic difference in parity: if the absolute value is less than 0.1, then the model can be considered fair.
- 2) Balanced odds difference: if the absolute value is less than 0.1, then the model can be considered fair.
- 3) The difference is in equal opportunities: if the absolute value is less than 0.1, then the model can be considered fair.
- 4) Demographic parity coefficient: the fairness of this indicator ranges from 0.8 to 1.25.

In our case, the difference in demographic parity (0.01), demographic parity ratio (0.99) does not indicate unfairness, while the difference in equalization of chances (0.34) indicates unfairness.

---

Impact of Mitigation: The application of the Demographic Parity constraint has resulted in a more balanced selection rate between genders. However, this has come with a trade-off in terms of accuracy and precision, especially for females.

Trade-off Between Fairness and Performance: The mitigation strategy seems to have prioritized demographic parity (equal selection rates), leading to a decrease in other performance metrics for females. This is a common challenge in fairness interventions, where improving fairness in one aspect can lead to decreases in performance metrics.

Fairness Considerations: Despite the trade-offs, the model is now less biased in terms of selection rate, which was one of the initial fairness concerns.

**Improvement in Fairness:** The intervention based on Equalized Odds has led to a reduction in both demographic and equalized odds disparities, indicating a more fair model.